GITHUB.COM/POOJAKIRA/MCP-AGENT-SECURITY-GATEWAY - PUBLIC, DEFAULT BRANCH MAIN.

Runtime Policy Enforcement at the AI Agent-to-Tool Boundary

Design and Validation of an Inline MCP/JSON-RPC Security Gateway

Evidence status: Refreshed against code snapshot f1f226feb0d527f27b90081806d7e9bd90ba0574 and successful GitHub Actions CI run 38502413461 on 2026-10-14 UTC. The quantitative claims below are tied to that code snapshot. Do not generalize benchmark or test results beyond their stated scope.

782
Tests passed
Python 3.12 CI; Python 3.10/3.11 also green

0
Tests failed
Cited CI run

81.93%
Statement coverage
5,657 statements, 1,022 missed

69
Prompt-injection collection entries
Named runtime collection; current main-branch implementation

9
Elastic rule records
detection_rules/elastic_rules.toml

21
Core SIEM tests
tests/test_siem.py

12
SIEM scenario tests
tests/test_siem_scenarios.py

Contribution

An application-layer inspection and authorization point for agent-generated MCP tools/call requests that combines JSON-RPC validation, normalization, trust/capability checks, heuristic content detection, policy evaluation, and tamper-evident audit telemetry before downstream tool execution.

Problem & Threat Model

LLM agents can turn untrusted natural-language context into structured tool calls with file, process, network, or privileged-service effects. The gateway protects the agent-to-tool boundary only when traffic is routed through it. Attackers may craft malicious tool arguments, malformed JSON-RPC, obfuscated prompt-injection strings, exfiltration payloads, or abusive request patterns.

Method

1. Parse and validate JSON-RPC, including duplicate-key and malformed-input rejection.
2. Normalize potentially obfuscated content before inspection.
3. Evaluate server trust, capabilities, and content/policy signals.
4. Return allow/block decisions using path-specific enforcement semantics.
5. Record hash-chained audit evidence and SIEM-oriented telemetry.

Reproducibility

The cited GitHub Actions run reports 782 passed and 81.93% statement coverage. Local results may differ with platform and dependencies; do not treat them as a guarantee for other commits.

Honest Boundaries

- No external population-level false-positive/false-negative rate is established.
- Synthetic correlation fixtures and local microbenchmarks are not production efficacy or SLO evidence.
- Enforcement semantics are not identical across every integration path.
- Only traffic routed through the gateway can be inspected or blocked.
- A successful CI suite is not evidence of a live production SOC deployment.

Evidence Sources & References

- GitHub Actions CI run 38502413461
- VERIFIED_METRICS.md
- evidence/mcp_replay_evidence.json
- detection_rules/elastic_rules.toml
- tests/test_siem.py
- tests/test_siem_scenarios.py